

# Tensorflow Basic (2/2)

# Optimization basic

- Training for prediction
  - $Y$  : a set target variable
  - $X$  : a set of feature vector to explain  $Y$
- Model selection to best capture the desired relation of  $X$  and  $Y$ 
  - Linear regression, logistic classification, deep neural network, ...
  - Our training data points will be used for tuning the model

# Optimization basic

- Defining a loss function
  - A good measure is required to evaluate the model's performance.
  - Measure the discrepancy or distance (loss) between the model's predictions and the observed targets.
  - The goal of optimization is to find the set of parameters (weight, biases, ...) of the model that minimize the distance.
  - Mean squared 
$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$
  - Cross entropy 
$$H(p, q) = - \sum_x p(x) \log q(x)$$

# Optimization basic

- Gradient descent: Finding the set of  $X$  that minimizes or  $F(X)$ .


- Derivative : 
$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

- Partial derivative,  $\nabla$

- $y = f(x_1, x_2, x_3, \dots, x_n)$ ,  $\nabla f = (f_{x_1}, f_{x_2}, f_{x_3}, \dots, f_{x_n})$

- Gradient descent

- $x' = x - \varepsilon * \nabla f$        $\varepsilon$ : learning rate



# Optimization basic

## \* Gradient descent optimizer

```
def _numerical_gradient_1d(f, x):  
    h = 1e-4 # 0.0001  
    grad = np.zeros_like(x)  
  
    for idx in range(x.size):  
        tmp_val = x[idx]  
        x[idx] = float(tmp_val) + h  
        fxh1 = f(x) # f(x+h)  
  
        x[idx] = tmp_val - h  
        fxh2 = f(x) # f(x-h)  
        grad[idx] = (fxh1 - fxh2) / (2*h)  
  
        x[idx] = tmp_val # 값 복원  
  
    return grad
```



f: loss function

x: parameters to tune

Iteratively update the set of parameters in a way that decrease the loss over time.

But the point we find could be a global minimum or a saddle point.


⇒ Second derivative



# Optimization basic

- Stochastic gradient descent
  - Instead of feeding the entire dataset to the algorithm for the computation of each iteration, a subset of the data is sampled (*mini-batch*) and fed.
  - Using smaller batches usually works faster and the smaller the size of the batch, the faster are the calculations. However, there is a trade-off in that small samples lead to lower hardware utilization and tend to have high variance, causing large fluctuations of the objective function.
  - Nevertheless, it turns out that some fluctuations are beneficial since they enable the set of parameters to jump to new and potentially better local minima.
  - Using a relatively small batches is preferred recently.

# Linear regression



Given:  $X = \{x_1, x_2, \dots, x_m\}$ ,  $Y = \{y_1, y_2, \dots, y_m\}$

Assumption:  $y_i \approx \hat{y}_i = a_1 * x_{i,1} + a_2 * x_{i,2} + \dots + a_m * x_{i,m} + b$

Goal: Find  $a_1, a_2, \dots, a_m, b$  that minimize

$\hat{Y} = \Phi * \theta$ , where  $\Phi = [x_1, x_2, \dots, x_m, 1]$ ,  $\theta = [a_1, a_2, \dots, a_m, b]^T$

$$\hat{Y} = \begin{bmatrix} [x_{1,1}, x_{1,2}, \dots, x_{1,m}, 1] \\ [x_{2,1}, x_{2,2}, \dots, x_{2,m}, 1] \\ \vdots \\ [x_{m,1}, x_{m,2}, \dots, x_{m,m}, 1] \end{bmatrix} * \begin{bmatrix} a_1 \\ a_2 \\ \vdots \\ a_m \\ b \end{bmatrix} = \Phi * \theta$$


# Linear regression with Tensorflow

## 1. Data generation

- $y$ : target variable
- $x$ : feature vector
- $y = x * W + b$

```
x_data = np.random.randn(2000, 3)
w_real = [0.3, 0.5, 0.1]
b_real = -0.2
noise = np.random.randn(1, 2000) * 0.1

temp = np.matmul(w_real, x_data.T) + b_real
y_data = temp + noise
```





# Linear regression with Tensorflow

## 2. Model definition

```
# training to predict
NUM_STEPS = 10

x = tf.placeholder(dtype=tf.float32, shape=[None, 3])
y_true = tf.placeholder(dtype=tf.float32, shape=None)

with tf.name_scope('inference') as scope:
    w = tf.Variable([[0,0,0]], dtype=tf.float32, name='weights')
    b = tf.Variable(0, dtype=tf.float32, name='bias')
    y_pred = tf.matmul(w, tf.transpose(x)) + b
```

## 3. Loss function definition

```
with tf.name_scope('mse') as scope:
    mse_loss = tf.reduce_mean(tf.square(y_true - y_pred))
```

```
with tf.name_scope('cross_entropy') as scope:
    cross_entropy_loss = tf.reduce_mean(tf.nn.sigmoid_cross_entropy_with_logits(labels=y_true, logits=y_pred))
```

# Linear regression with Tensorflow

## 4. Training

```
with tf.name_scope('train') as scope:  
    learning_rate = 0.5  
    optimizer = tf.train.GradientDescentOptimizer(learning_rate)  
    train = optimizer.minimize(mse_loss)
```

```
with tf.Session() as sess:  
    sess.run(tf.global_variables_initializer())  
    for step in range(NUM_STEPS):  
        sess.run(train, feed_dict={x:x_data, y_true:y_data})  
        if step % 5 == 0:  
            print step, sess.run([w, b])  
  
    print 10, sess.run([w, b])
```

## 5. Result

```
0 [array([[0.3105729 , 0.49970657, 0.08790597]], dtype=float32), -0.21092714]  
5 [array([[0.29926383, 0.49876568, 0.09481297]], dtype=float32), -0.1982763]  
10 [array([[0.29926383, 0.49876568, 0.09481298]], dtype=float32), -0.1982763]
```