

Lecture 15: Red-Black Trees

Hyungon Moon

Acknowledgment: The content of this file is based on the slides of the textbook as well as the slides provided by Prof. Won-Ki Jeong.

Red-Black Trees

- Another self-balancing binary search tree
- Guarantee $O(\log n)$ insertion, search, delete
- Definition
 - Binary search tree that every node is colored either red or black
 - Leaf nodes do not contain data
 - External nodes
 - i.e., every node has either 0 or 2 children.


Red-Black Tree Properties

- The root and all external nodes are black
- No root-to-external-node path has two consecutive red nodes
 - (=) Red node must have two black children
- All root-to-external-node paths have the same number of black nodes

Red-Black Tree Properties (Pointer)

- Pointers from an internal node to an external node are black
- No root-to-external-node path has two consecutive red pointers
- All root-to-external-node paths have the same number of black pointers

Example Red-Black Tree



Properties

- If P and Q are two root-to-external-node paths in a red-black tree, then


$$Length_P \leq 2Length_Q$$

$$\text{i.e., } Length_{Longest} \leq 2Length_{Shortest}$$

- Shortest path : B-B-B-....-B
- Longest path : B-R-B-R...-B
- Number of B must be same for all paths by definition

Properties

- Rank : # of black edges on any path from a node to any external node ($= L_{Shortest}$)




Properties: $h = O(\log n)$

- h : height, r : rank of the root, n : # of nodes
- $h \leq 2r$
 - Discussed earlier.
- $n \geq 2^r - 1$
 - When all nodes are black.
- $h \leq 2\log_2(n + 1) \Rightarrow h = O(\log n)$


Inserting

- Just insert and then color
- How to color a new node?
 - If the tree was empty, new node is root so assign black
 - If the tree was not empty, assign black causes increase one black node in the path : NO!
 - b/c violate same # of black nodes for all paths, difficult resolve
 - If the tree was not empty, assign red may cause two consecutive red nodes in the path : OK!
 - Can be resolved by rotation and color flips


Possible consecutive reds



LL(RR)b




LL(RR)b




XYr

XYr \Rightarrow color flip




- Flip color of pp, gp, d and pointers of gp
- Flip color d to black is ok b/c gp is also flipped
- Reapply transformation to gp by $p = gp$

XYr \Rightarrow color flip




- The number of blacks in each path is preserved.
- No two consecutive reds (in this part).
- Need to flip recursively.

LLb \Rightarrow rotate and flip



- Same as LL rotation of AVL tree
- Flip color of pp and gp after rotation
- No need to check parent; root color is not changed

LRb \Rightarrow rotate twice and flip




- Same as LR rotation of AVL tree
- Flip color of p and gp
- RRb and RLb are symmetric

Insert Example

- Insert 1
 - Root




- Insert 2
 - Red




Insert Example

- Insert 3
 - RRb




Insert Example

- Insert 4
– RRR




Insert Example

- Insert 5
 - RRB




Delete

- Similar to insert, but more complicated
- Delete black will violate red-black property
 - Path passing through deleted node will have less number of black nodes
 - Double black pointer




delete p



white circle : any color can be placed


Delete

- Case 0 : w (sibling) and its children are black
 - Make w red
 - If z is black, then move up double black pointer. $x = z, z = z \rightarrow \text{parent}$ and restart (below z is ok)




Delete

- Case 1 : z is black and w is red
 - Left-rotate at z and exchange colors of z & w
 - Go to case 2, 3, or 4 for subtree of z




Delete

- Case 2 : w and its two children are black
 - Make w red
 - If z is red, then make z black and remove double pointer. Done.




Delete

- Case 3 : w and its right child are black while its left child is red
 - Right-rotate at w and exchange colors of w and its left child. Go to case 4.




Delete

- Case 4 : w is black and its right child is red
 - Left rotate at z, exchange colors of z & w
 - Remove double black pointer, change R to black
 - Done



Delete Workflow

- At most 3 rotations are needed
- Color exchange may propagate $\log n$ times



Discussion

- Red-Black trees use color as balancing information instead of height as in AVL trees
- Insertion/deletion may cause a perturbation (if two consecutive red nodes exist)
- Perturbation is either
 - resolved locally (rotations), or
 - propagated to a higher level in the tree by recoloring (color flip)
- $O(1)$ for a rotation or $O(\log n)$ color flips
- Total time: $O(\log n)$

Questions?