

Course Overview

CSE251: System Programming
1st Lecture, Feb. 25, 2019

Instructor:

Hyungon Moon

All slides for this course are from:

<https://www.cs.cmu.edu/afs/cs/academic/class/15213-f15/www/schedule.html>

Overview

- Course theme
- Five realities
- Course logistics

Course Theme:

Abstraction Is Good But Don't Forget Reality

■ Most CSE courses emphasize abstraction

- Abstract data types
- Asymptotic analysis

■ These abstractions have limits

- Especially in the presence of bugs
- Need to understand details of underlying implementations

■ Useful outcomes from taking 251

- Become more effective programmers
 - Able to find and eliminate bugs efficiently
 - Able to understand and tune for program performance
- Prepare for later “systems” classes in CSE
 - Compilers, Operating Systems, Networks, Computer Architecture, etc.

Great Reality #1:


Ints are not Integers, Floats are not Reals

■ Example 1: Is $x^2 \geq 0$?

■ Float's: Yes!

■ Int's:

- $40000 * 40000 \rightarrow 1600000000$
- $50000 * 50000 \rightarrow ??$



■ Example 2: Is $(x + y) + z = x + (y + z)$?

■ Unsigned & Signed Int's: Yes!

■ Float's:

- $(1e20 + -1e20) + 3.14 \rightarrow 3.14$
- $1e20 + (-1e20 + 3.14) \rightarrow ??$

Computer Arithmetic

■ Does not generate random values

- Arithmetic operations have important mathematical properties

■ Cannot assume all “usual” mathematical properties

- Due to finiteness of representations
- Integer operations satisfy “ring” properties
 - Commutativity, associativity, distributivity
- Floating point operations satisfy “ordering” properties
 - Monotonicity, values of signs

■ Observation

- Need to understand which abstractions apply in which contexts
- Important issues for compiler writers and serious application programmers

Great Reality #2:

You've Got to Know Assembly

- **Chances are, you'll never write programs in assembly**
 - Compilers are much better & more patient than you are
- **But: Understanding assembly is key to machine-level execution model**
 - Behavior of programs in presence of bugs
 - High-level language models break down
 - Tuning program performance
 - Understand optimizations done / not done by the compiler
 - Understanding sources of program inefficiency
 - Implementing system software
 - Compiler has machine code as target
 - Operating systems must manage process state
 - Creating / fighting malware
 - x86 assembly is the language of choice!

Great Reality #3: Memory Matters

Random Access Memory Is an Unphysical Abstraction

■ Memory is not unbounded

- It must be allocated and managed
- Many applications are memory dominated

■ Memory referencing bugs especially pernicious

- Effects are distant in both time and space

■ Memory performance is not uniform

- Cache and virtual memory effects can greatly affect program performance
- Adapting program to characteristics of memory system can lead to major speed improvements

Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;  
  
double fun(int i) {  
    volatile struct_t s;  
    s.d = 3.14;  
    s.a[i] = 1073741824; /* Possibly out of bounds */  
    return s.d;  
}
```

fun(0)	→	3.14
fun(1)	→	3.14
fun(2)	→	3.13999998664856
fun(3)	→	2.000000061035156
fun(4)	→	3.14
fun(6)	→	Segmentation fault

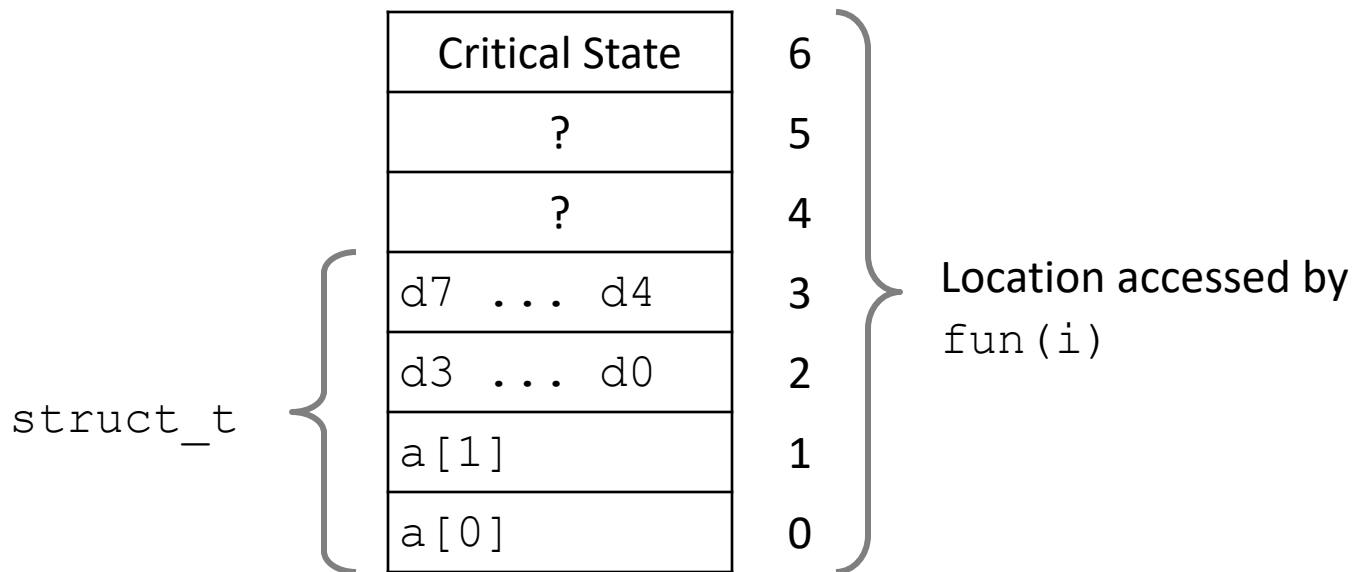
- Result is system specific

Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;
```

fun(0)	→	3.14
fun(1)	→	3.14
fun(2)	→	3.1399998664856
fun(3)	→	2.00000061035156
fun(4)	→	3.14
fun(6)	→	Segmentation fault

Explanation:



Memory Referencing Errors

■ C and C++ do not provide any memory protection

- Out of bounds array references
- Invalid pointer values
- Abuses of malloc/free

■ Can lead to nasty bugs

- Whether or not bug has any effect depends on system and compiler
- Action at a distance
 - Corrupted object logically unrelated to one being accessed
 - Effect of bug may be first observed long after it is generated

■ How can I deal with this?

- Program in Java, Ruby, Python, ML, ...
- Understand what possible interactions may occur
- Use or develop tools to detect referencing errors (e.g. Valgrind, asan)

Great Reality #4: There's more to performance than asymptotic complexity

- **Constant factors matter too!**
- **And even exact op count does not predict performance**
 - Easily see 10:1 performance range depending on how code written
 - Must optimize at multiple levels: algorithm, data representations, procedures, and loops
- **Must understand system to optimize performance**
 - How programs compiled and executed
 - How to measure program performance and identify bottlenecks
 - How to improve performance without destroying code modularity and generality

Memory System Performance Example

```
void copyij(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

4.3ms

2.0 GHz Intel Core i7 Haswell

```
void copyji(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

81.8ms

- Hierarchical memory organization
- Performance depends on access patterns
 - Including how step through multi-dimensional array

Great Reality #5:

Computers do more than execute programs

- **They need to get data in and out**
 - I/O system critical to program reliability and performance
- **They communicate with each other over networks**
 - Many system-level issues arise in presence of network
 - Concurrent operations by autonomous processes
 - Coping with unreliable media
 - Cross platform compatibility
 - Complex performance issues

Course Perspective

- Most Systems Courses are Builder-Centric
 - Computer Architecture
 - Design pipelined processor in Verilog
 - Operating Systems
 - Implement sample portions of operating system
 - Compilers
 - Write compiler for simple language
 - Networking
 - Implement and simulate network protocols

Course Perspective (Cont.)

■ Our Course is Programmer-Centric

- Purpose is to show that by knowing more about the underlying system, one can be more effective as a programmer
- Enable you to
 - Write programs that are more reliable and efficient
 - Incorporate features that require hooks into OS
 - E.g., concurrency, signal handlers
- Cover material in this course that you won't see elsewhere
- Not just a course for dedicated hackers
 - **We bring out the hidden hacker in everyone!**


Logistics: Teaching staff

■ Instructor

- Hyungon Moon (hyungon@unist.ac.kr)
- Office hour: by email
 - Best timing to email: ~ 8am in the morning.
- Office: 106-701-8

■ TAs

- Sehoon Kim (sshhhee@unist.ac.kr) @ 106-605
- TBD1
- TBD2
- TBD3



Grading

■ General rule

- 5%: Class participation
- 20%: Midterm
- 25%: Final
- 50%: Lab

■ Special rule

- Not submitting two or more lab will make your grade C+ or lower.
- We follow university and school policies about attendance and cheating.
 - e.g., Missing 8 or more classes will make your grade F.
- Class participation: online discussion, in-class participation.
 - **Doesn't** include the attendance.

The *official* ECE rule on the Academic Integrity

■ On the 1st violation

- **Zero grade on the item** involved (e.g., homework, midterm, etc.).
- **Lower the final grade by at least one letter grade** (e.g., B+ --> C+).
- Attain a **“signed” personal letter from the student** stating this will not happen again, and he/she is well aware of consequence if it does.
- **Provide a written report** of the student and violation **to School Head and ECE education committee.**

■ On the 2nd violation

- **Give an F on the course.**
- **Share the identity of the student** with the entire faculty.
- **Report to the University Student Scholarship Counseling Committee** for further disciplinary action.

Textbook

- Randal E. Bryant and David R. O'Hallaron,
 - *Computer Systems: A Programmer's Perspective*, **Third Edition** (CS:APP3e), Pearson, 2016
 - Reference: <http://csapp.cs.cmu.edu>
 - This book really matters for the course!
 - How to solve labs
 - Practice problems typical of exam problems

Course Components

■ Lectures

- Higher level concepts

■ Recitations

- Applied concepts, important tools and skills for labs, clarification of lectures, exam coverage

■ Labs (7+1)

- The heart of the course
- 1-2 weeks each
- Provide in-depth understanding of an aspect of systems
- Programming and measurement

■ Exams (midterm + final)

- Test your understanding of concepts & mathematical principles

Getting Help

- Class Gitlab: <https://class.unicss.org/cse251-2019-spring>
- Accounts:
 - Registered students will get password reset email by 2/26. Get registered!
- General questions:
 - Create/check issues on the master repository.
 - <https://class.unicss.org/cse251-2019-spring>
- Private questions:
 - You are not likely to have these.
 - Send an email or create an issue on your submission repository.
- Blackboard or Piazza
 - We won't be using Blackboard or Piazza for the course

Policies

■ Work groups

- You must work alone on all lab assignments

■ Handins

- Labs due at 11:59pm on each deadline
- Electronic handins using the class Gitlab.
- The Lab 0 will cover the basics of git.

Facilities

- Labs will be tested on the uni server
 - `linux> ssh uni06.unist.ac.kr`
- Getting help with the uni server
 - Please create an issue on the class Gitlab.

Timeliness

- No late submissions allowed.
- Depending on the labs, you may get earlier feedback:
 - You will have a chance to revise your submission, only if you submit it early.
- ***Start early***
 - It could take quite a lot of time: you should start early to finish and submit it.

Other Rules of the class room

- Laptops: permitted
- Electronic communications: **forbidden**
 - No email, instant messaging, cell phone calls, etc
- Presence in lectures, recitations: voluntary, recommended
- No recordings of ANY KIND

Programs and Data

■ Topics

- Becoming familiar with git, the de facto standard version control system.

■ Assignments

- L0 (gitlab): Trying out the lab infra and becoming familiar with Linux shell.

Programs and Data

■ Topics

- Bits operations, arithmetic, assembly language programs
- Representation of C control and data structures
- Includes aspects of architecture and compilers

■ Assignments

- L1 (datalab): Manipulating bits
- L2 (bomblab): Defusing a binary bomb
- L3 (attacklab): The basics of code injection attacks

The Memory Hierarchy

■ Topics

- Memory technology, memory hierarchy, caches, disks, locality
- Includes aspects of architecture and OS

■ Assignments

- L4 (cachelab): Building a cache simulator and optimizing for locality.
 - Learn how to exploit locality in your programs.

Exceptional Control Flow

■ Topics

- Hardware exceptions, processes, process control, Unix signals, nonlocal jumps
- Includes aspects of compilers, OS, and architecture

■ Assignments

- L5 (tshlab): Writing your own Unix shell.
 - A first introduction to concurrency

Virtual Memory

■ Topics

- Virtual memory, address translation, dynamic storage allocation
- Includes aspects of architecture and OS

■ Assignments

- L6 (malloclab): Writing your own malloc package
 - Get a real feel for systems-level programming

Networking, and Concurrency

■ Topics

- High level and low-level I/O, network programming
- Internet services, Web servers
- concurrency, concurrent server design, threads
- I/O multiplexing with select
- Includes aspects of networking, OS, and architecture

■ Assignments

- L7 (proxylab): Writing your own Web proxy
 - Learn network programming and more about concurrency and synchronization.

Lab Rationale

- Each lab has a well-defined goal such as solving a puzzle or winning a contest
- Doing the lab should result in new skills and concepts
- We try to use competition in a fun and healthy way
 - Set a reasonable threshold for full credit

Gitlab showcase

- There are plenty of resources that you can follow.
- The steps you are going to follow:
 - Create a password and login
 - Read the syllabus (README.md file on the master repository)
 - Create and register your ssh key
 - <https://docs.gitlab.com/ee/ssh/>

Welcome
and Enjoy!