

Shell Lab: Exceptional Control Flow

CSE251, Spring 2019

Recitation 5: Wed, Apr 24th, 2019

Changmin Yi

ulistar93@unist.ac.kr

*Reference : CMU 15-213: Intro to Computer Systems Fall 2015
Recitation 9 - Celeste Neary*

Shell Lab

- Shell Lab is out!
- Due Thursday, May. 2, 11:59PM
 - **Please don't be late!**
- The grader will run every 6am/noon/6pm on 4/24 – 5/2
- The final grader will run 00:05 5/3

Shell Lab

```
Makefile
myint*
myint.c
myspin*
myspin.c
mysplit*
mysplit.c
mystop*
mystop.c
README.md
README.txt
sdriver.pl*
shlab.pdf
trace01.txt
trace02.txt
trace03.txt
trace04.txt
trace05.txt
trace06.txt
trace07.txt
trace08.txt
trace09.txt
trace10.txt
trace11.txt
trace12.txt
trace13.txt
trace14.txt
trace15.txt
trace16.txt
tsh*
tsh.c
tshref*
tshref.out
```

- There are a lot of files, but there is an only one file you should implement.
tsh.c
- You will implement...
 - `eval`: Main routine that parses and interprets the command line. [70 lines]
 - `builtin_cmd`: Recognizes and interprets the built-in commands: `quit`, `fg`, `bg`, and `jobs`. [25 lines]
 - `do_bgfg`: Implements the `bg` and `fg` built-in commands. [50 lines]
 - `waitfg`: Waits for a foreground job to complete. [20 lines]
 - `sigchld_handler`: Catches `SIGCHLD` signals. 80 lines]
 - `sigint_handler`: Catches `SIGINT` (`ctrl-c`) signals. [15 lines]
 - `sigtstp_handler`: Catches `SIGTSTP` (`ctrl-z`) signals. [15 lines]

Shell Lab Testing

- You can run your shell by
 - `./tsh`
(you should 'make' first whatever you run)
 - `tsh` is a shell as like Unix shell

Unix

```
[cs20141494@uni06 shlab]$ /bin/ps
PID TTY      TIME CMD
 6986 pts/35    00:00:00 ps
16199 pts/35    00:00:00 bash
```

tsh

```
[cs20141494@uni06 shlab]$ ./tsh
tsh> /bin/ps
PID TTY      TIME CMD
 6987 pts/35    00:00:00 tsh
 6988 pts/35    00:00:00 ps
16199 pts/35    00:00:00 bash
tsh>
tsh> [cs20141494@uni06 shlab]$
```

← run ./tsh

exit tsh
"ctrl + d"
= "^d"

Shell Lab Testing

- You can run your shell by
 - `./tsh`
- The possible commands are
 - `/bin/l`s
 - `/bin/p`s
 - `/bin/e`cho
 - `jobs`
 - `bg`
 - `fg`
 - `kill`
 - `quit`
- Please do not try in `tsh>`
 - `vi`, `emacs`, `more`, `less`

Shell Lab Testing

- And you can test your shell by

- sdriver.pl

```
unix> ./sdriver.pl -h
Usage: sdriver.pl [-hv] -t <trace> -s <shellprog> -a <args>
Options:
  -h                Print this message
  -v                Be more verbose
  -t <trace>        Trace file
  -s <shell>        Shell program to test
  -a <args>         Shell arguments
  -g                Generate output for autograder
```

- For example

```
unix> ./sdriver.pl -t trace01.txt -s ./tsh -a "-p"
```

or

```
unix> make test01
```

(both are same)

Shell Lab Testing

- And also there is reference answer shell !
 - So your `./tsh` should be work same as `./tshref`
- The reference answer is
 - `tshref.out`

```
make[1]: Entering directory `/afs/cs.cmu.edu/project/ics/im/labs/shlab/src'
./sdriver.pl -t trace01.txt -s ./tsh -a "-p"
#
# trace01.txt - Properly terminate on EOF.
#
./sdriver.pl -t trace02.txt -s ./tsh -a "-p"
#
# trace02.txt - Process builtin quit command.
#
./sdriver.pl -t trace03.txt -s ./tsh -a "-p"
#
# trace03.txt - Run a foreground job.
#
tsh> quit
./sdriver.pl -t trace04.txt -s ./tsh -a "-p"
#
# trace04.txt - Run a background job.
#
tsh> ./myspin 1 &
[1] (26252) ./myspin 1 &
```

Shell Lab

- While doing the implementation,
- You need to consider the hazards that already told in lecture
 - Race conditions
 - Hard to debug so start early (and think carefully)
 - Reaping zombies
 - Race conditions
 - Handling signals correctly
 - Waiting for foreground job
 - Think carefully about what the right way to do this is
- **Hints** in the instructions will be helpful when you work
I recommend reading it carefully

Processes

- Four basic States
 - Running
 - Executing instructions on the CPU
 - Number bounded by number of CPU cores
 - Runnable
 - Waiting to be running
 - Blocked
 - Waiting for an event, maybe input from STDIN
 - Not runnable
 - Zombie
 - Terminated, not yet reaped

Processes

- Four basic process control function families:
 - `fork()`
 - `exec()`
 - And other variants such as `execve()`
 - `exit()`
 - `wait()`
 - And variants like `waitpid()`
- Standard on all UNIX-based systems

Signals

- A *signal* is a small message that notifies a process that an event of some type has occurred in the system
 - akin to exceptions and interrupts (asynchronous)
 - sent from the kernel (sometimes at the request of another process) to a process
 - signal type is identified by small integer ID's (1-30)
 - only information in a signal is its ID and the fact that it arrived

<i>ID</i>	<i>Name</i>	<i>Default Action</i>	<i>Corresponding Event</i>
2	SIGINT	Terminate	Interrupt (e.g., ctrl-c from keyboard)
9	SIGKILL	Terminate	Kill program (cannot override or ignore)
11	SIGSEGV	Terminate & Dump	Segmentation violation
14	SIGALRM	Terminate	Timer signal
17	SIGCHLD	Ignore	Child stopped or terminated

Signals

- Kernel *sends* (delivers) a signal to a *destination process* by updating some state in the context of the destination process

- Kernel sends a signal for one of the following reasons:
 - Kernel has detected a system event such as Ctrl-C (SIGINT), divide-by-zero (SIGFPE), or the termination of a child process (SIGCHLD)
 - Another program called the `kill()` function
 - The user used a `kill` utility

Signals

- A destination process *receives* a signal when it is forced by the kernel to react in some way to the delivery of the signal
- Receiving a signal is non-queuing
 - There is only one bit in the context per signal
 - Receiving 1 or 300 SIGINTs looks the same to the process
- Signals are received at a context switch
- Three possible ways to react:
 - *Ignore* the signal (do nothing)
 - *Terminate* the process (with optional core dump)
 - *Catch* the signal by executing a user-level function called *signal handler*
 - Akin to a hardware exception handler being called in response to an asynchronous interrupt

Signals

- A destination process *receives* a signal when it is forced by the kernel to react in some way to the delivery of the signal
- Blocking signals
 - Sometimes code needs to run through a section that can't be interrupted
 - Implemented with `sigprocmask()`
- Waiting for signals
 - Sometimes, we want to pause execution until we get a specific signal
 - Implemented with `sigsuspend()`
 - > What different with `wait()` and `waitpid()`?
- Can't modify behavior of SIGKILL and SIGSTOP

Signals

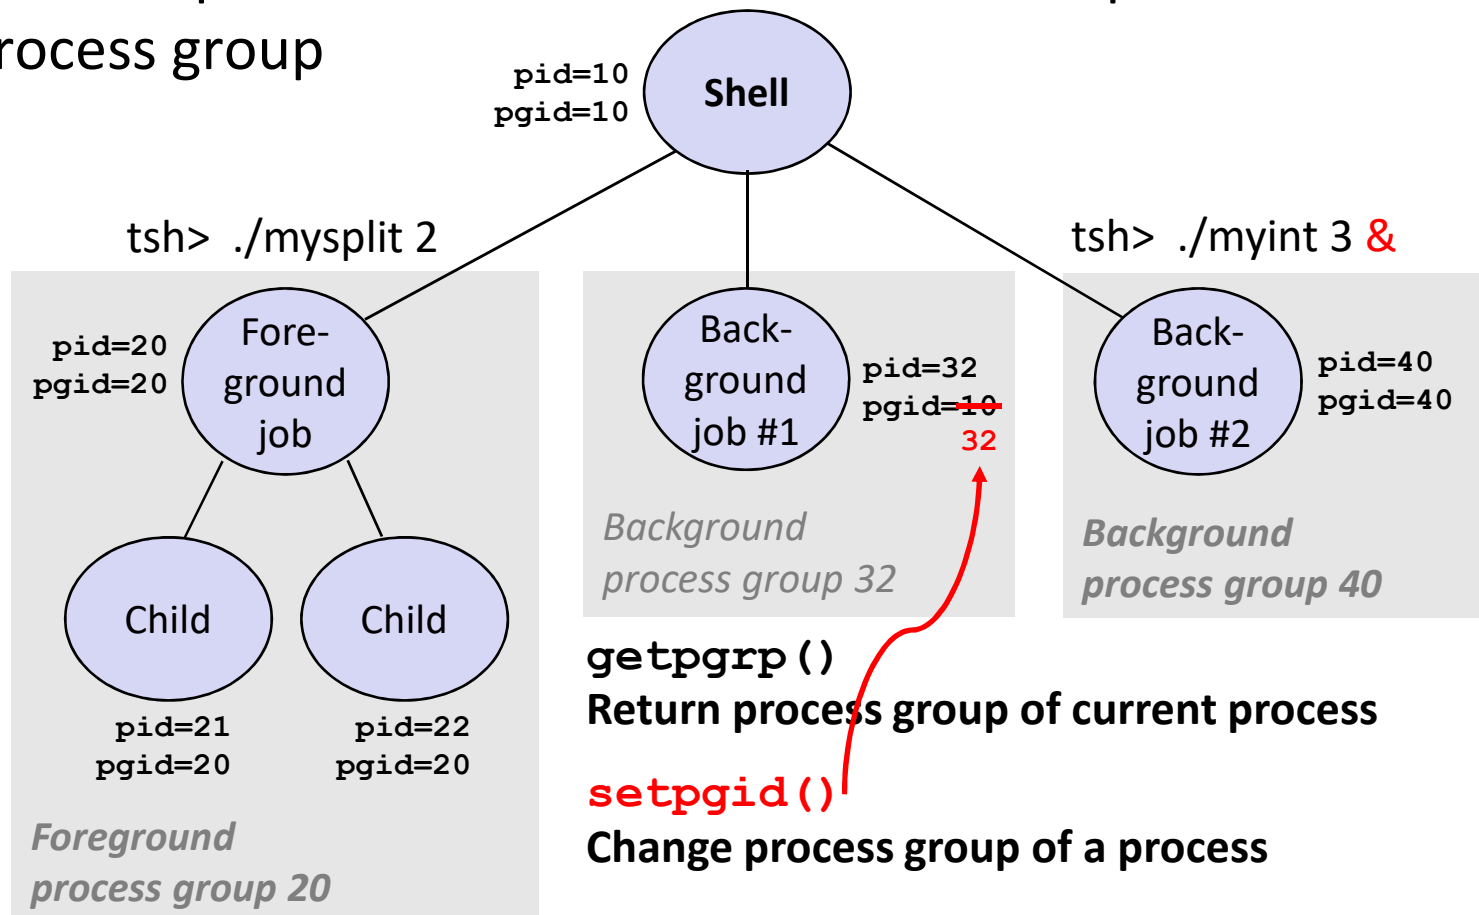
- Signal handlers
 - Can be installed to run when a signal is received
 - The form is `void handler(int signum){ ... }`
 - **Separate** flow of control in the same process
 - Resumes normal flow of control upon returning
 - Can be called **anytime** when the appropriate signal is fired

Signals

- `int sigsuspend(const sigset_t *mask)`
 - Can't use `wait()` twice – use `sigsuspend`!
 - Temporarily replaces the signal mask of the calling process with the mask given
 - Suspends the process until delivery of a signal whose action is to invoke a signal handler or terminate a process
 - Returns if the signal is caught
 - Signal mask restored to the previous state
 - Use `sigaddset()`, `sigemptyset()`, etc. to create the mask

Signal Examples

- Every process belongs to exactly one process group
- Process groups can be used to distribute signals easily
- A forked process becomes a member of the parent's process group



Signal Examples

```
// sigchld handler installed

pid_t child_pid = fork();

if (child_pid == 0){
    /* child comes here */

    execve(.....);
}
else{

    add_job(child_pid);

}
```

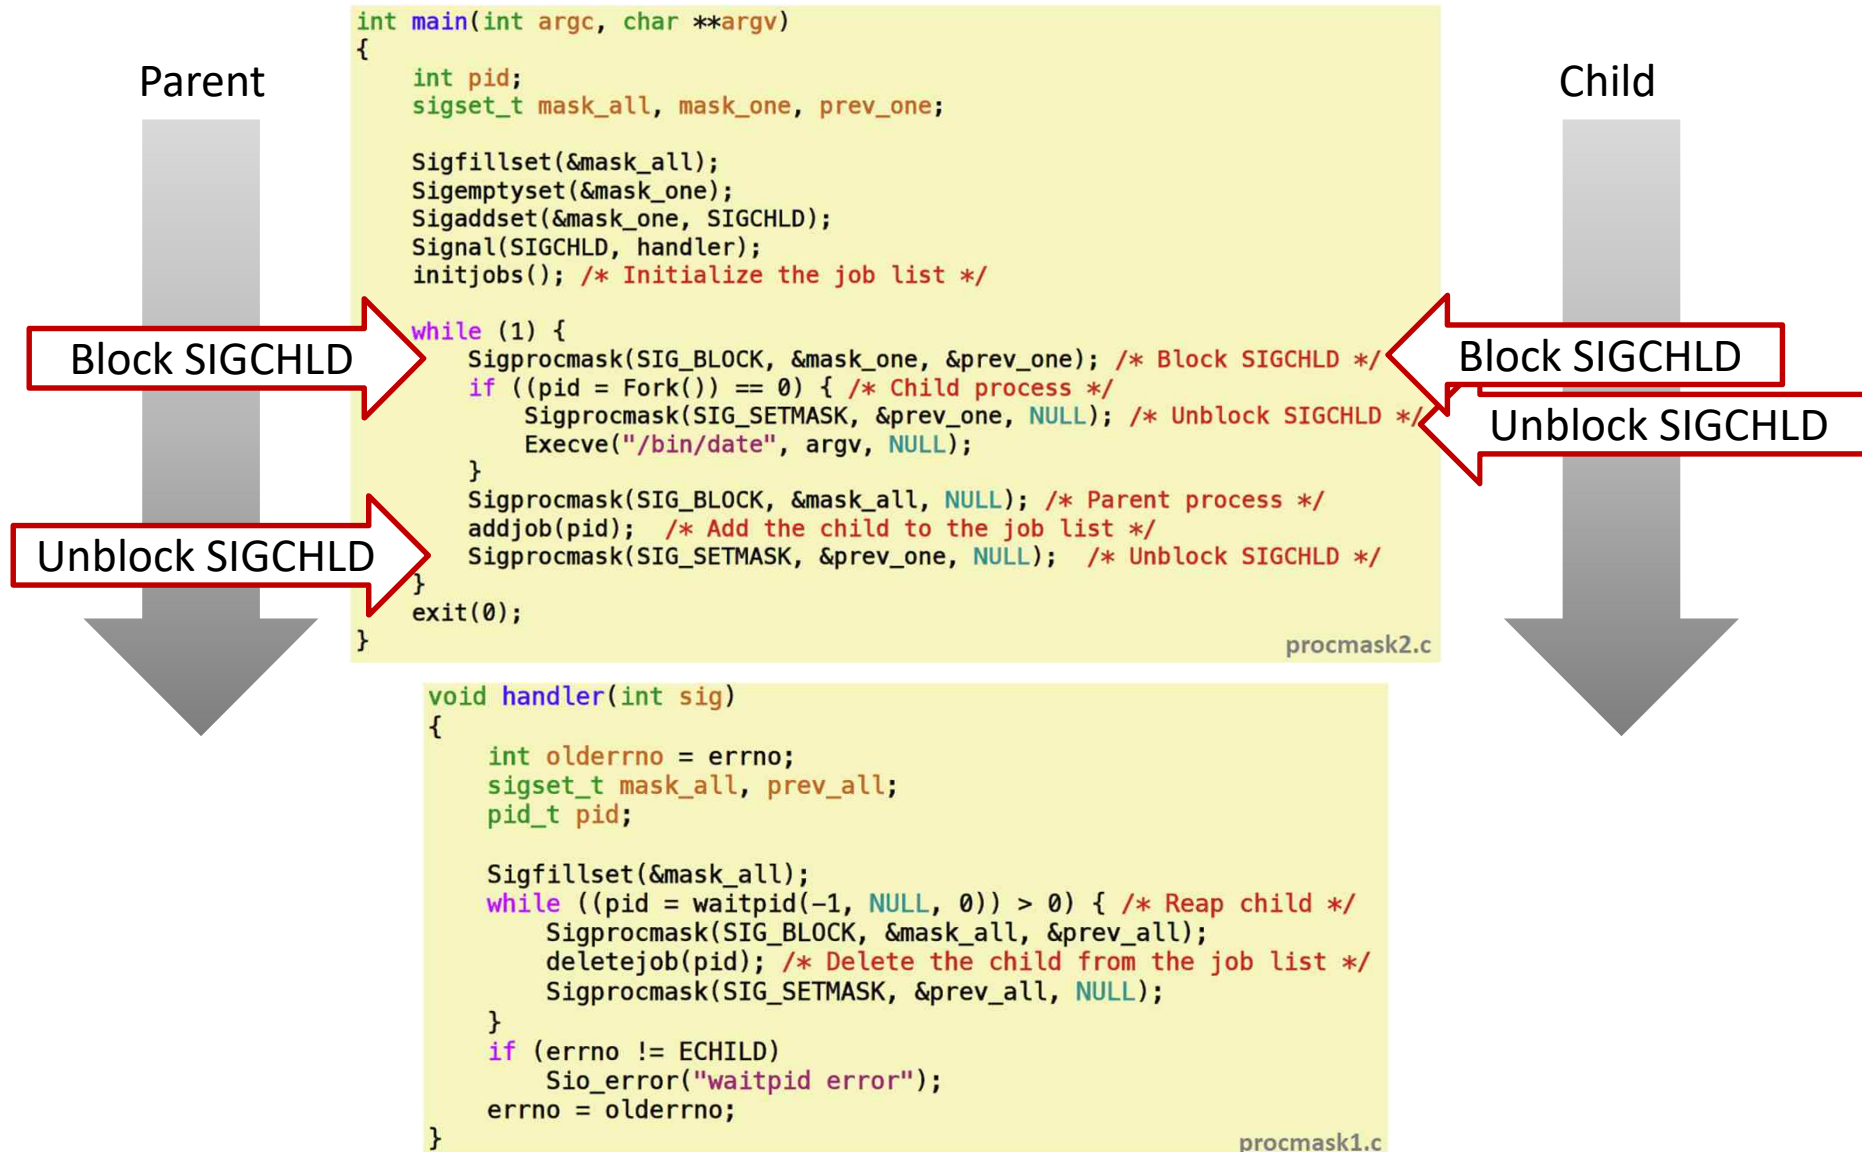
```
void sigchld_handler(int signum)
{
    int status;

    pid_t child_pid =
        waitpid(-1, &status, WNOHANG);

    if (WIFEXITED(status))
        remove_job(child_pid);
}
```

- Does add_job or remove_job() come first?
- Where can we block signals in this code to guarantee correct execution?

Corrected Shell Program without Race



Async-Signal-Safety

- Function is *async-signal-safe* if either reentrant (e.g., all variables stored on stack frame, CS:APP3e 12.7.2) or non-interruptible by signals.
- Posix guarantees 117 functions to be async-signal-safe
 - Source: “man 7 signal”
 - Popular functions on the list:
 - `_exit`, `write`, `wait`, `waitpid`, `sleep`, `kill`
 - Popular functions that are **not** on the list:
 - `printf`, `sprintf`, `malloc`, `exit`
 - Unfortunate fact: `write` is the only async-signal-safe output function

Safely Generating Formatted Output

- Use the reentrant SIO (Safe I/O library) from `csapp.c` in your handlers.

- `ssize_t sio_puts(char s[]) /* Put string */`
- `ssize_t sio_putl(long v) /* Put long */`
- `void sio_error(char s[]) /* Put msg & exit */`

```
void sigint_handler(int sig) /* Safe SIGINT handler */
{
    Sio_puts("So you think you can stop the bomb with ctrl-
c, do you?\n");
    sleep(2);
    Sio_puts("Well...\n");
    sleep(1);
    Sio_puts("OK. :-)\n");
    _exit(0);
}
```

sigintsafe.c